

Suffix Array Algorithm



Join a Google user study to shape the future of our products and get a gift for your input.

SIGN UP

asked 3 years ago

viewed 14999 times

active 1 year ago

 40

After quite a bit of reading, I have figured out what a suffix array and LCP array represents.

 35

Suffix array: Represents the `_lexicographic` rank of each suffix of an array.

LCP array : Contains the maximum length prefix match between two consecutive suffixes, after they are *sorted lexicographically*.

I have been trying hard to understand since a couple of days , how exactly the **suffix array and LCP algorithm works**.

Here is the code , which is taken from [Codeforces](#):

```
/*
Suffix array O(n lg^2 n)
LCP table O(n)
*/
#include <cstdio>
#include <algorithm>
#include <cstring>

using namespace std;

#define REP(i, n) for (int i = 0; i < (int)(n); ++i)

namespace SuffixArray
{
```

Join a Google user study to shape the future of our products and get a gift for your input.



SIGN UP

Looking for a job?

C++ Market Data Feeds Developer

Stevens Capital Management LP  Radnor, PA

 **RELOCATION**  **VISA SPONSORSHIP**

c++

linux

Product Manager, Internal Systems

```

const int MAXN = 1 << 21;
char * S;
int N, gap;
int sa[MAXN], pos[MAXN], tmp[MAXN], lcp[MAXN];

bool sufCmp(int i, int j)
{
    if (pos[i] != pos[j])
        return pos[i] < pos[j];
    i += gap;
    j += gap;
    return (i < N && j < N) ? pos[i] < pos[j] : i > j;
}

void buildSA()
{
    N = strlen(S);
    REP(i, N) sa[i] = i, pos[i] = S[i];
    for (gap = 1;; gap *= 2)
    {
        sort(sa, sa + N, sufCmp);
        REP(i, N - 1) tmp[i + 1] = tmp[i] + sufCmp(sa[i], sa[i + 1]);
        REP(i, N) pos[sa[i]] = tmp[i];
        if (tmp[N - 1] == N - 1) break;
    }
}

```

I cannot, just cannot get through how this algorithm works. I tried working on an example using pencil and paper, and wrote through the steps involved, but lost link in between as its too complicated, for me at least.

Any help regarding explanation, using an example maybe, is highly appreciated.

c++ algorithm data-structures suffix-array

share improve this question

edited Jul 20 '13 at 13:02

asked Jul 20 '13 at 11:21



Spandan

739 1 15 30

2 #define REP – don't do this. Nobody will use such code. If you don't like the classical for loop (which is very understandable!) then you can use for (i in rep(n)) instead, if you implement rep accordingly (here's a rough guide: github.com/klmr/cpp11-range) – Konrad Rudolph Jul 20 '13 at 12:25

4 @KonradRudolph: ok..i will do that.perhaps,now something for my questn! – Spandan Jul 20 '13 at 12:35

Product manager, Internal Systems

Stack Overflow New York, NY

\$100,000 - \$135,000

product

user-experience

Linked

7

Implementation of string pattern matching using Suffix Array and LCP(-LR)

10

Suffix array nlogn creation

2

How does this code for obtaining LCP from a Suffix Array work?

0

How to sort a suffix array?

Related

12

How to sort array suffixes in block sorting

34

What's the current state-of-the-art suffix array construction algorithm?

766

Ukkonen's suffix tree algorithm in plain English?

16539

Why is it faster to process a sorted array than an unsorted array?

I recently had to implement a suffix array for a class. Searching in it is just a binary search with "is the pattern lexicographically smaller than or equal to the suffix?" as the search condition. I didn't understand how the code to build the suffix array works, though - which we were actually given. The only suggestion I can offer is: if you really have the time to do this, you could have a look at the [original paper](#) where the suffix array was first proposed. – [G. Bach](#) Jul 20 '13 at 13:20

Using the string "mississippi" the `buildLCP` accesses `s` at index `N` which is out of bounds unless the string is required to be null terminated. Is it? – [Jon Harrop](#) Sep 15 '15 at 0:50

[add a comment](#)

2 Answers

active

oldest

votes



83



Overview

This is an $O(n \log n)$ algorithm for suffix array construction (or rather, it would be, if instead of `::sort` a 2-pass bucket sort had been used).

It works by first sorting the 2-grams^(*), then the 4-grams, then the 8-grams, and so forth, of the original string `s`, so in the i -th iteration, we sort the 2^i -grams. There can obviously be no more than $\log_2(n)$ such iterations, and the trick is that sorting the 2^i -grams in the i -th step is facilitated by making sure that each comparison of two 2^i -grams is done in $O(1)$ time (rather than $O(2^i)$ time).

How does it do this? Well, **in the first iteration** it sorts the 2-grams (aka bigrams), and then performs what is called *lexicographic renaming*. This means it creates a new array (of length `n`) that stores, for each bigram, its *rank* in the bigram sorting.

Example for lexicographic renaming: Say we have a **sorted** list of some bigrams

`{'ab', 'ab', 'ca', 'cd', 'cd', 'ea'}`. We then assign *ranks* (i.e. lexicographic names) by going from left to right, starting with rank 0 and incrementing the rank whenever we encounter a *new* bigram changes. So the ranks we assign are as follows:

```
ab : 0
ab : 0    [no change to previous]
ca : 1    [increment because different from previous]
cd : 2    [increment because different from previous]
cd : 2    [no change to previous]
ea : 3    [increment because different from previous]
```

These ranks are known as *lexicographic names*.

4

[Minimum Lexicographic Rotation Using Suffix Array](#)

3

[suffix array using manber myers algorithm](#)

1

[Suffix array DC3 algorithm](#)

1

[Understanding implementation of DC3/Skew algorithm to create Suffix Array linear time](#)

0

[Suffix Array Construction Algorithm](#)

6

[Using suffix array algorithm for Burrows Wheeler transform](#)

Hot Network Questions

[Output the sign](#)

[Is it bad form to write mysterious proofs without explaining what one intends to do?](#)

[Calculate the prime factors](#)

[Why didn't Darth Vader tell Sidious that Obi-Wan disappeared instead of dying?](#)

[Is this an allowed step for working with infinite sequences?](#)

[Are there statistics of airports by flight cancellation rate?](#)

Now, **in the next iteration**, we sort 4-grams. This involves a lot of comparisons between different 4-grams. How do we compare two 4-grams? Well, we could compare them character by character. That would be up to 4 operations per comparison. But instead, we compare them by *looking up* the ranks of the two bigrams contained in them, using the rank table generated in the previous steps. That rank represents the lexicographic rank from the previous 2-gram sort, so if for any given 4-gram, its first 2-gram has a higher rank than the first 2-gram of another 4-gram, then it must be lexicographically greater *somewhere in the first two characters*. Hence, if for two 4-grams the rank of the first 2-gram is identical, they must be identical in the *first two characters*. In other words, *two look-ups* in the rank table are sufficient to compare all 4 characters of the two 4-grams.

After sorting, we create new lexicographic names again, this time for the 4-grams.

In the third iteration, we need to sort by 8-grams. Again, two look-ups in the lexicographic rank table from the previous step are sufficient to compare all 8 characters of two given 8-grams.

And so forth. Each iteration `i` has two steps:

1. Sorting by 2^i -grams, using the lexicographic names from the previous iteration to enable comparisons in 2 steps (i.e. $O(1)$ time) each
2. Creating new lexicographic names



















We repeat this until all 2^i -grams are different. If that happens, we are done. How do we know if all are different? Well, the lexicographic names are an increasing sequence of integers, starting with 0. So if the highest lexicographic name generated in an iteration is the same as `n-1`, then each 2^i -gram must have been given its own, distinct lexicographic name.

Implementation

Now let's look at the code to confirm all of this. The variables used are as follows: `sa[]` is the suffix array we are building. `pos[]` is the rank lookup-table (i.e. it contains the lexicographic names), specifically, `pos[k]` contains the lexicographic name of the `k`-th m-gram of the previous step. `tmp[]` is an auxiliary array used to help create `pos[]`.

I'll give further explanations between the code lines:

```
void buildSA()
{
    N = strlen(S);
```

-  Could aliens colonize Earth without realizing humans are people too?
-  Word for "exploding error"?
-  Explain it to me like I'm a physics grad: Greenhouse Effect
-  Is it possible to see packet before encryption?
-  My js file is not loaded in Magento 2
-  Change \baselineskip to match other fontsize than currently used
-  When hiking, why is the right of way given to people going up?
-  Stopping text from being italic in shaded environment
-  Why did Blofeld tell Hans to keep the keys to spaceship's self-destruct?
-  Does a Ghana citizen need \$3000.00 in hand to travel to the USA?
-  Is password-based AES encryption secure at all?
-  My office wants infinite branch merges as policy; what other options do we have?
-  How can I convince players not to offload a seemingly useless weapon?
-  What signature prevents me from spending others' coins?
-  Having a list of paths, how can I filter out subdirectories of previously mentioned paths?
-  Terminal - Is the built in zip command lossy?
-  On my Canon Rebel (EOS T6i), is it possible to re-program the default settings of some of the shooting modes? I want flash-off to be the default
-  Betti numbers as characteristic numbers?

```

/* This is a loop that initializes sa[] and pos[].
   For sa[] we assume the order the suffixes have
   in the given string. For pos[] we set the lexicographic
   rank of each 1-gram using the characters themselves.
   That makes sense, right? */
REP(i, N) sa[i] = i, pos[i] = S[i];

/* Gap is the length of the m-gram in each step, divided by 2.
   We start with 2-grams, so gap is 1 initially. It then increases
   to 2, 4, 8 and so on. */
for (gap = 1;; gap *= 2)
{
    /* We sort by (gap*2)-grams: */
    sort(sa, sa + N, sufCmp);

    /* We compute the lexicographic rank of each m-gram
       that we have sorted above. Notice how the rank is computed
       by comparing each n-gram at position i with its
       neighbor at i+1. If they are identical, the comparison
       yields 0, so the rank does not increase. Otherwise the
       comparison yields 1, so the rank increases by 1. */
    REP(i, N - 1) tmp[i + 1] = tmp[i] + sufCmp(sa[i], sa[i + 1]);

    /* tmp contains the rank by position. Now we map this
       into pos, so that in the next step we can look it
       up per m-gram, rather than by position. */
    REP(i, N) pos[sa[i]] = tmp[i];

    /* If the largest lexicographic name generated is
       n-1, we are finished, because this means all
       m-grams must have been different. */
    if (tmp[N - 1] == N - 1) break;
}
}

```

About the comparison function

The function `sufCmp` is used to compare two $(2 \cdot \text{gap})$ -grams lexicographically. So in the first iteration it compares bigrams, in the second iteration 4-grams, then 8-grams and so on. This is controlled by `gap`, which is a global variable.

A naive implementation of `sufCmp` would be this:

```

bool sufCmp(int i, int j)
{
    int pos_i = sa[i];

```

```

int pos_j = sa[j];

int end_i = pos_i + 2*gap;
int end_j = pos_j + 2*gap;
if (end_i > N)
    end_i = N;
if (end_j > N)
    end_j = N;

while (i < end_i && j < end_j)
{
    if (S[pos_i] != S[pos_j])
        return S[pos_i] < S[pos_j];
    pos_i += 1;
    pos_j += 1;
}
return (pos_i < N && pos_j < N) ? S[pos_i] < S[pos_j] : pos_i > pos_j;
}

```

This would compare the $(2*gap)$ -gram at the beginning of the i -th suffix `pos_i:=sa[i]` with the one found at the beginning of the j -th suffix `pos_j:=sa[j]`. And it would compare them character by character, i.e. comparing `S[pos_i]` with `S[pos_j]`, then `S[pos_i+1]` with `S[pos_j+1]` and so on. It continues as long as the characters are identical. Once they differ, it returns 1 if the character in the i -th suffix is smaller than the one in the j -th suffix, 0 otherwise. (Note that `return a<b` in a function returning `int` means you return 1 if the condition is true, and 0 if it is false.)

The complicated looking condition in the return-statement deals with the case that one of the $(2*gap)$ -grams is located at the end of the string. In this case either `pos_i` or `pos_j` will reach `N` before all $(2*gap)$ characters have been compared, even if all characters up to that point are identical. It will then return 1 if the i -th suffix is at the end, and 0 if the j -th suffix is at the end. This is correct because if all characters are identical, the *shorter* one is lexicographically smaller. If `pos_i` has reached the end, the i -th suffix must be shorter than the j -th suffix.

Clearly, this naive implementation is $O(gap)$, i.e. its complexity is linear in the length of the $(2*gap)$ -grams. The function used in your code, however, uses the lexicographic names to bring this down to $O(1)$ (specifically, down to a maximum of two comparisons):

```

bool sufCmp(int i, int j)
{
    if (pos[i] != pos[j])
        return pos[i] < pos[j];
    i += gap;
    j += gap;
    return (i < N && j < N) ? pos[i] < pos[j] : i > j;
}

```

```
}
```

As you can see, instead of looking up individual characters `s[i]` and `s[j]`, we check the lexicographic rank of the *i*-th and *j*-th suffix. Lexicographic ranks were computed in the previous iteration for gap-grams. So, if `pos[i] < pos[j]`, then the *i*-th suffix `sa[i]` must start with a gap-gram that is lexicographically smaller than the gap-gram at the beginning of `sa[j]`. In other words, simply by looking up `pos[i]` and `pos[j]` and comparing them, we have compared the first *gap* characters of the two suffixes.

If the ranks are identical, we continue by comparing `pos[i+gap]` with `pos[j+gap]`. This is the same as comparing the next *gap* characters of the ($2 \cdot \text{gap}$)-grams, i.e. the *second half*. If the ranks are identical again, the two ($2 \cdot \text{gap}$)-grams are identical, so we return 0. Otherwise we return 1 if the *i*-th suffix is smaller than the *j*-th suffix, 0 otherwise.

Example

The following example illustrates how the algorithm operates, and demonstrates in particular the role of the lexicographic names in the sorting algorithm.

The string we want to sort is `abcxabcd`. It takes three iterations to generate the suffix array for this. In each iteration, I'll show `s` (the string), `sa` (the current state of the suffix array) and `tmp` and `pos`, which represent the lexicographic names.

First, we initialize:

```
S   abcxabcd
sa  01234567
pos abcxabcd
```

Note how the lexicographic names, which initially represent the lexicographic rank of unigrams, are simply identical to the characters (i.e. the unigrams) themselves.

First iteration:

Sorting `sa`, using bigrams as sorting criterion:

```
sa  04156273
```

The first two suffixes are 0 and 4 because those are the positions of bigram 'ab'. Then 1 and 5 (positions of bigram 'bc'), then 6 (bigram 'cd'), then 2 (bigram 'cx'). then 7 (incomplete bigram 'd'), then

3 (bigram 'xa'). Clearly, the positions correspond to the order, based solely on character bigrams.

Generating the lexicographic names:

```
tmp 00112345
```

As described, lexicographic names are assigned as increasing integers. The first two suffixes (both starting with bigram 'ab') get 0, the next two (both starting with bigram 'bc') get 1, then 2, 3, 4, 5 (each a different bigram).

Finally, we map this according to the positions in `sa`, to get `pos`:

```
sa 04156273
tmp 00112345
pos 01350124
```

(The way `pos` is generated is this: Go through `sa` from left to right, and use the entry to define the index in `pos`. Use the corresponding entry in `tmp` to define the value for that index. So `pos[0]:=0`, `pos[4]:=0`, `pos[1]:=1`, `pos[5]:=1`, `pos[6]:=2`, and so on. The index comes from `sa`, the value from `tmp`.)

Second iteration:

We sort `sa` again, and again we look at bigrams from `pos` (which each represents a sequence of *two* bigrams of the original string).

```
sa 04516273
```

Notice how the position of 1 5 have switched compared to the previous version of `sa`. It used to be 15, now it is 51. This is because the bigram at `pos[1]` and the bigram at `pos[5]` used to be identical (both `bc`) in during the previous iteration, but now the bigram at `pos[5]` is `12`, while the bigram at `pos[1]` is `13`. So position `5` comes *before* position `1`. This is due to the fact that the lexicographic names now each represent bigrams of the original string: `pos[5]` represents `bc` and `pos[6]` represents `cd`. So, together they represent `bcd`, while `pos[1]` represents `bc` and `pos[2]` represents `cx`, so together they represent `bcx`, which is indeed lexicographically greater than `bcd`.

Again, we generate lexicographic names by screening the current version of `sa` from left to right and comparing the corresponding bigrams in `pos`:

```
tmp 00123456
```


The first two entries are still identical (both 0), because the corresponding bigrams in `pos` are both `01`. The rest is an strictly increasing sequence of integers, because all other bigrams in `pos` are each unique.

We perform the mapping to the new `pos` as before (taking indices from `sa` and values from `tmp`):

```
sa  04516273
tmp 00123456
pos 02460135
```

Third iteration:

We sort `sa` again, taking bigrams of `pos` (as always), which now each represents a sequence of 4 bigrams of the original string.

```
sa  40516273
```

You'll notice that now the first two entries have switched positions: `04` has become `40`. This is because the bigram at `pos[0]` is `02` while the one at `pos[4]` is `01`, the latter obviously being lexicographically smaller. The deep reason is that these two represent `abcx` and `abcd`, respectively.

Generating lexicographic names yields:

```
tmp 01234567
```

They are all different, i.e. the highest one is `7`, which is `n-1`. So, we are done, because are sorting is now based on m-grams that are all different. Even if we continued, the sorting order would not change.

Improvement suggestion

The algorithm used to sort the 2^i -grams in each iteration appears to be the built-in `sort` (or `std::sort`). This means it's a comparison sort, which takes $O(n \log n)$ time in the worst case, *in each iteration*. Since there are $\log n$ iterations in the worst case, this makes it a $O(n (\log n)^2)$ -time algorithm. However, the sorting could be performed using two passes of bucket sort, since the keys we use for the sort comparison (i.e. the lexicographic names of the previous step), form an increasing integer sequence. So this could be improved to an actual $O(n \log n)$ -time algorithm for suffix sorting.

Remark

I believe this is the original algorithm for suffix array construction that was suggested in the 1992-paper by Manber and Myers ([link on Google Scholar](#); it should be the first hit, and it may have a link to a PDF there). This (at the same time, but independently of a paper by Gonnet and Baeza-Yates) was what introduced suffix arrays (also known as pat arrays at the time) as a data structure interesting for further study.

Modern algorithms for suffix array construction are $O(n)$, so the above is no longer the best algorithm available (at least not in terms of theoretical, worst-case complexity).

Footnotes

(*) By *2-gram* I mean a sequence of two *consecutive* characters of the original string. For example, when $S=abcde$ is the string, then ab , bc , cd , de are the 2-grams of S . Similarly, $abcd$ and bcd are the 4-grams. Generally, an m -gram (for a positive integer m) is a sequence of m consecutive characters. 1-grams are also called unigrams, 2-grams are called bigrams, 3-grams are called trigrams. Some people continue with tetragrams, pentagrams and so on.

Note that the suffix of S that starts and position i , is an $(n-i)$ -gram of S . Also, every m -gram (for any m) is a prefix of one of the suffixes of S . Therefore, sorting m -grams (for an m as large as possible) can be the first step towards sorting suffixes.

share improve this answer

edited Oct 6 '15 at 6:45



Mark VY

308 1 11

answered Jul 20 '13 at 15:05



jogojapan

41k 7 61 88

1 @jogojapan: When i asked this questn, i just hoped , that you read it .Only then i would have got such a comprehensive answer, given your illustrious history for answering quuestions related to Suffix tree/array . Thanks a lot sir . Perhaps,i was hoping ,only if u had time , you could write a bit bout d construction of LCp array too.thanks. – [Spandan](#) Jul 20 '13 at 18:21

1 @G.Bach Uhm. Yes, it is. The variable name “k” changes of course, but you’re the only one on this page using it. – [Konrad Rudolph](#) Jul 20 '13 at 19:28

1 @jogojapan: i'm hoping you'll add the LCP array construction too,in some time. – [Spandan](#) Jul 21 '13 at 13:25

1 @DhruvMullick Thanks for asking the question about the LCP code [here](#). I hope I answered the question about [k](#) (which called [l](#) there) at the other question. Let's discuss there if not. – [jogojapan](#) Oct 18 '14 at 9:54

1 @Spandan I'm sorry this took forever, but now somebody asked again about the LCP algorithm, and I hope I was able to finally answer it here: [stackoverflow.com/questions/26428636/...](#) – [jogojapan](#) Oct 18 '14 at 9:56


[show 14 more comments](#)



Monetize your app with over 1 million Google advertisers

Learn more



 3 As you've visited Codeforces for this, I believe you're familiar with Topcoder too :) Please check the link below and it has some more links too. I've seen many people learning Suffix array from these links only.

 [TopCoder Link Here](#)

And a piece of gold for Programming Contestants. I learned Suffix array from [this paper](#)

[share](#) [improve this answer](#)

[edited Mar 6 '14 at 5:52](#)

[answered Jul 20 '13 at 19:22](#)



Fallen

2,416 1 8 26

thanks...will look into it – [Spandan](#) Jul 20 '13 at 19:36

the second link (paper) doesn't work anymore, could you please upload it or give us another link? – [fersarr](#) Mar 5 '14 at 2:25

is it this one ? [arxiv.org/pdf/0912.0807v1.pdf](#) – [fersarr](#) Mar 5 '14 at 2:27

1 @fersarr: Link updated. Please check that link now :) [stanford.edu/class/cs97si/suffix-array.pdf](#) – [Fallen](#) Mar 6 '14 at 5:53

[add a comment](#)

Your Answer

B B B B B B B B B B B B

Sign up or [log in](#)

Post as a guest



Sign up using Google



Sign up using Facebook



Sign up using Email and Password

Name

Email

Post Your Answer

By posting your answer, you agree to the [privacy policy](#) and [terms of service](#).

Not the answer you're looking for? Browse other questions tagged [c++](#) [algorithm](#) [data-structures](#)

[suffix-array](#) or [ask your own question](#).

[about us](#) [tour](#) [help](#) [blog](#) [chat](#) [data](#) [legal](#) [privacy policy](#) [work here](#) [advertising info](#) [mobile](#) [contact us](#) [feedback](#)

TECHNOLOGY

Stack Overflow	Software Engineering	Database Administrators	Code Review
Server Fault	Unix & Linux	Drupal Answers	Magento
Super User	Ask Different (Apple)	SharePoint	Signal Processing
Web Applications	WordPress Development	User Experience	Raspberry Pi
Ask Ubuntu	Geographic Information Systems	Mathematica	Programming Puzzles & Code Golf
Webmasters	Electrical Engineering	Salesforce	more (7)
Game Development	Android Enthusiasts	ExpressionEngine® Answers	
TeX - LaTeX	Information Security	Cryptography	

LIFE / ARTS

Photography	Academia
Science Fiction & Fantasy	more (8)
Graphic Design	
Movies & TV	
Music: Practice & Theory	
Seasoned Advice (cooking)	
Home Improvement	
Personal Finance & Money	

CULTURE / RECREATION

English Language & Usage	Bicycles
Skeptics	Role-playing Games
Mi Yodeya (Judaism)	Anime & Manga
Travel	Motor Vehicle Maintenance & Repair
Christianity	more (17)
English Language Learners	
Japanese Language	
Arqade (gaming)	

SCIENCE

MathOverflow
Mathematics
Cross Validated (stats)
Theoretical Computer Science
Physics
Chemistry
Biology
Computer Science

OTHER

Philosophy
more (3)
Meta Stack Exchange
Stack Apps
Area 51
Stack Overflow Talent